

Liver Disease Prediction

Content

This data set contains 416 liver patient records and 167 non liver patient records collected from North East of Andhra Pradesh, India. The "Dataset" column is a class label used to divide groups into liver patient (liver disease) or not (no disease). This data set contains 441 male patient records and 142 female patient records.

Any patient whose age exceeded 89 is listed as being of age "90".

Columns:

Age of the patient

Gender of the patient

Total Bilirubin

Direct Bilirubin

Alkaline Phosphatase

Alamine Aminotransferase

Aspartate Aminotransferase

Total Protiens

Albumin

Albumin and Globulin Ratio

Dataset: field used to split the data into two sets (patient with liver disease, or no disease)

Importing Libraries

Importing Libraries:

```
import pandas as pd
```

```
import numpy as np
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
# for displaying all feature from dataset:
```

```
pd.pandas.set_option('display.max_columns', None)
```

```
Read The DataSet
```

```
# Reading Dataset:
```

```
dataset = pd.read_csv("https://raw.githubusercontent.com/IBM-EPBL/IBM-Project-7380-1658854276/main/project/Dataset/indian_liver_patient.csv")
```

```
# Top 5 records:
```

```
dataset.head()
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	Albumin	Albumin_and_Globulin_Ratio	Dataset
0	65	Female	0.7	0.1	187	16	18	6.8	3.3	0.90	1
1	62	Male	10.9	5.5	699	64	100	7.5	3.2	0.74	1
2	62	Male	7.3	4.1	490	60	68	7.0	3.3	0.89	1
3	58	Male	1.0	0.4	182	14	20	6.8	3.4	1.00	1
4	72	Male	3.9	2.0	195	27	59	7.3	2.4	0.40	1

```
# Last 5 records:
```

```
dataset.tail()
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	Albumin	Albumin_and_Globulin_Ratio	Dataset
578	60	Male	0.5	0.1	500	20	34	5.9	1.6	0.37	2
579	40	Male	0.6	0.1	98	35	31	6.0	3.2	1.10	1
580	52	Male	0.8	0.2	245	48	49	6.4	3.2	1.00	1

581	31	Male	1.3	0.5	184	29	32	6.8	3.4	1.00	1
582	38	Male	1.0	0.3	216	21	24	7.3	4.4	1.50	2

Shape of dataset:

dataset.shape

(583, 11)

Cheaking Missing (NaN) Values

Cheaking Missing (NaN) Values:

dataset.isnull().sum()

Age 0

Gender 0

Total_Bilirubin 0

Direct_Bilirubin 0

Alkaline_Phosphotase 0

Alamine_Aminotransferase 0

Aspartate_Aminotransferase 0

Total_Protiens 0

Albumin 0

Albumin_and_Globulin_Ratio 4

Dataset 0

dtype: int64

'Albumin_and_Globulin_Ratio' feature contain 4 NaN values.

Mean & Median of "Albumin_and_Globulin_Ratio" feature:

print(dataset['Albumin_and_Globulin_Ratio'].median())

print(dataset['Albumin_and_Globulin_Ratio'].mean())

Age	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	Albumin	Albumin_and_Globulin_Ratio
count	583.000000	583.000000	583.000000	583.000000	583.000000	583.000000	583.000000	583.000000
	583.000000	583.000000	583.000000	583.000000				

mean	44.746141 6.483190	3.298799 3.141852	1.486106 0.946947	290.576329 1.286449	80.713551	109.910806
std	16.189833 1.085451	6.209522 0.795519	2.808498 0.318495	242.937989 0.452490	182.620356	288.918529
min	4.000000 2.700000	0.400000 0.900000	0.100000 0.300000	63.000000 1.000000	10.000000	10.000000
25%	33.000000 5.800000	0.800000 2.600000	0.200000 0.700000	175.500000 1.000000	23.000000	25.000000
50%	45.000000 6.600000	1.000000 3.100000	0.300000 0.930000	208.000000 1.000000	35.000000	42.000000
75%	58.000000 7.200000	2.600000 3.800000	1.300000 1.100000	298.000000 2.000000	60.500000	87.000000
max	90.000000 9.600000	75.000000 5.500000	19.700000 2.800000	2110.000000 2.000000	2000.000000	4929.000000

Data Visualization

Target feature:

```
print("Liver Disease Patients :", dataset['Dataset'].value_counts()[1])
```

```
print("Non Liver Disease Patients :", dataset['Dataset'].value_counts()[2])
```

Visualization:

```
sns.countplot(dataset['Dataset'])
```

```
plt.show()
```

Liver Disease Patients : 416

Non Liver Disease Patients : 167

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Histogram of Age:

```
plt.figure(figsize=(8,5))
```

```
sns.histplot(dataset['Age'], kde=True)
```

```
plt.title('Age', fontsize=20)
```

```
plt.show()
```

```
dataset.head()
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	Albumin	Albumin_and_Globulin_Ratio	Dataset
0	65	Female	0.7	0.1	187	16	18	6.8	3.3	0.90	1
1	62	Male	10.9	5.5	699	64	100	7.5	3.2	0.74	1
2	62	Male	7.3	4.1	490	60	68	7.0	3.3	0.89	1
3	58	Male	1.0	0.4	182	14	20	6.8	3.4	1.00	1
4	72	Male	3.9	2.0	195	27	59	7.3	2.4	0.40	1

Gender feature:

```
print("Total Male :", dataset['Gender'].value_counts()[0])
```

```
print("Total Female :", dataset['Gender'].value_counts()[1])
```

Visualization:

```
sns.countplot(dataset['Gender'])
```

```
plt.show()
```

Total Male : 441

Total Female : 142

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Printing How many Unique values present in each feature:

for feature in dataset.columns:

```
print(feature,":", len(dataset[feature].unique()))
```

Age : 72

Gender : 2

Total_Bilirubin : 113

Direct_Bilirubin : 80

Alkaline_Phosphotase : 263

Alamine_Aminotransferase : 152

Aspartate_Aminotransferase : 177

Total_Protiens : 58

Albumin : 40

Albumin_and_Globulin_Ratio : 69

Dataset : 2

Label Encoding

```
dataset['Gender'] = np.where(dataset['Gender']=='Male', 1,0)
```

```
dataset.head()
```

Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	Albumin	Albumin_and_Globulin_Ratio	Dataset	
0	65	0	0.7	0.1	187	16	18	6.8	3.3	0.90	1
1	62	1	10.9	5.5	699	64	100	7.5	3.2	0.74	1

2	62	1	7.3	4.1	490	60	68	7.0	3.3	0.89	1
3	58	1	1.0	0.4	182	14	20	6.8	3.4	1.00	1
4	72	1	3.9	2.0	195	27	59	7.3	2.4	0.40	1

Correlation using Heatmap:

```
plt.figure(figsize=(12,8))

sns.heatmap(dataset.corr(), annot=True, cmap='YlGnBu')

plt.show()
```

There is Multi-Collinearity found on our dataset.

dataset.columns

```
Index(['Age', 'Gender', 'Total_Bilirubin', 'Direct_Bilirubin',
       'Alkaline_Phosphotase', 'Alamine_Aminotransferase',
       'Aspartate_Aminotransferase', 'Total_Protiens', 'Albumin',
       'Albumin_and_Globulin_Ratio', 'Dataset'],
      dtype='object')
```

Multicollinearity between 'Total_Bilirubin' and 'Direct_Bilirubin' is 0.87%

Multicollinearity between 'Alamine_Aminotransferase' and 'Aspartate_Aminotransferase' is 0.79%

Multicollinearity between 'Total_Protiens' and 'Albumin' is 0.78%

Multicollinearity between 'Albumin' and 'Albumin_and_Globulin_Ratio' is 0.69%

Usually we drop that feature which has above 0.85% multicollinearity between two independent feature. Here we have only 'Total_Bilirubin' and 'Direct_Bilirubin' feature which has 0.87% multicollinearity. So we drop one of the feature from them and other independent feature has less multicollinearity, less than 0.80% So we keep that feature.

Dropping 'Direct_Bilirubin' feature:

```
dataset = dataset.drop('Direct_Bilirubin', axis=1)

dataset.columns
```



```
Index(['Age', 'Gender', 'Total_Bilirubin', 'Alkaline_Phosphotase',  
      'Alamine_Aminotransferase', 'Aspartate_Aminotransferase',  
      'Total_Protiens', 'Albumin', 'Albumin_and_Globulin_Ratio', 'Dataset'],  
      dtype='object')
```

```
sns.distplot(dataset['Albumin'])
```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
# Calculate the boundaries of Total_Protiens feature which differentiates the outliers:
```

```
uppper_boundary=dataset['Total_Protiens'].mean() + 3* dataset['Total_Protiens'].std()
```

```
lower_boundary=dataset['Total_Protiens'].mean() - 3* dataset['Total_Protiens'].std()
```

```
print(dataset['Total_Protiens'].mean())
```

```
print(lower_boundary)
```

```
print(uppper_boundary)
```

```
6.483190394511149
```

```
3.2268359424407516
```

```
9.739544846581545
```

```
##### Calculate the boundaries of Albumin feature which differentiates the outliers:
```

```
uppper_boundary=dataset['Albumin'].mean() + 3* dataset['Albumin'].std()
```

```
lower_boundary=dataset['Albumin'].mean() - 3* dataset['Albumin'].std()
```

```
print(dataset['Albumin'].mean())
```

```
print(lower_boundary)
```

```
print(upper_boundary)
```

```
3.141852487135506
```

```
0.7552960692434296
```

```
5.528408905027582
```

```
# Lets compute the Interquantile range of Total_Bilirubin feature to calculate the boundaries:
```

```
IQR = dataset.Total_Bilirubin.quantile(0.75)-dataset.Total_Bilirubin.quantile(0.25)
```

```
# Extreme outliers
```

```
lower_bridge = dataset['Total_Bilirubin'].quantile(0.25) - (IQR*3)
```

```
upper_bridge = dataset['Total_Bilirubin'].quantile(0.75) + (IQR*3)
```

```
print(lower_bridge)
```

```
print(upper_bridge)
```

```
# if value greater than upper bridge, we replace that value with upper_bridge value:
```

```
dataset.loc[dataset['Total_Bilirubin'] >= upper_bridge, 'Total_Bilirubin'] = upper_bridge
```

```
-4.6000000000000005
```

```
8.0
```

```
# Lets compute the Interquantile range of Alkaline_Phosphotase feature to calculate the boundaries:
```

```
IQR = dataset.Alkaline_Phosphotase.quantile(0.75) - dataset.Alkaline_Phosphotase.quantile(0.25)
```

```
# Extreme outliers
```

```
lower_bridge = dataset['Alkaline_Phosphotase'].quantile(0.25) - (IQR*3)
```

```
upper_bridge = dataset['Alkaline_Phosphotase'].quantile(0.75) + (IQR*3)
```

```
print(lower_bridge)
```

```
print(upper_bridge)
```

```
# if value greater than upper bridge, we replace that value with upper_bridge value:
```

```
dataset.loc[dataset['Alkaline_Phosphotase'] >= upper_bridge, 'Alkaline_Phosphotase'] = upper_bridge
```

```
-192.0
```

```
665.5
```

```
# Lets compute the Interquantile range of Alamine_Aminotransferase feature to calculate the boundaries:
```

```
IQR = dataset.Alamine_Aminotransferase.quantile(0.75) -  
dataset.Alamine_Aminotransferase.quantile(0.25)
```

```
# Extreme outliers
```

```
lower_bridge = dataset['Alamine_Aminotransferase'].quantile(0.25) - (IQR*3)
```

```
upper_bridge = dataset['Alamine_Aminotransferase'].quantile(0.75) + (IQR*3)
```

```
print(lower_bridge)
```

```
print(upper_bridge)
```

```
# if value greater than upper bridge, we replace that value with upper_bridge value:
```

```
dataset.loc[dataset['Alamine_Aminotransferase'] >= upper_bridge, 'Alamine_Aminotransferase'] =  
upper_bridge
```

```
-89.5
```

```
173.0
```

```
# Lets compute the Interquantile range of Aspartate_Aminotransferase feature to calculate the boundaries:
```

```
IQR = dataset.Aspartate_Aminotransferase.quantile(0.75) -  
dataset.Aspartate_Aminotransferase.quantile(0.25)
```

```
# Extreme outliers
```

```
lower_bridge = dataset['Aspartate_Aminotransferase'].quantile(0.25) - (IQR*3)
```

```
upper_bridge = dataset['Aspartate_Aminotransferase'].quantile(0.75) + (IQR*3)
```

```
print(lower_bridge)
```

```
print(upper_bridge)
```

```
# if value greater than upper bridge, we replace that value with upper_bridge value:
```

```
dataset.loc[dataset['Aspartate_Aminotransferase'] >= upper_bridge, 'Aspartate_Aminotransferase'] =  
upper_bridge
```

```
-161.0
```

```
273.0
```

```
# Lets compute the Interquantile range of Albumin_and_Globulin_Ratio feature to calculate the  
boundaries
```

```
IQR = dataset.Albumin_and_Globulin_Ratio.quantile(0.75) -  
dataset.Albumin_and_Globulin_Ratio.quantile(0.25)
```

```
# Extreme outliers
```

```
lower_bridge = dataset['Albumin_and_Globulin_Ratio'].quantile(0.25) - (IQR*3)
```

```
upper_bridge = dataset['Albumin_and_Globulin_Ratio'].quantile(0.75) + (IQR*3)
```

```
print(lower_bridge)
```

```
print(upper_bridge)
```

if value greater than upper bridge, we replace that value with upper_bridge value:

```
dataset.loc[dataset['Albumin_and_Globulin_Ratio'] >= upper_bridge, 'Albumin_and_Globulin_Ratio'] =  
upper_bridge
```

-0.50000000000000004

2.3000000000000007

Top 5 records:

```
dataset.head()
```

	Age	Gender	Total_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	Albumin	Albumin_and_Globulin_Ratio
0	65	0	0.7	187.0	16	18	6.8	3.3	0.90
1	62	1	8.0	665.5	64	100	7.5	3.2	0.74
2	62	1	7.3	490.0	60	68	7.0	3.3	0.89
3	58	1	1.0	182.0	14	20	6.8	3.4	1.00
4	72	1	3.9	195.0	27	59	7.3	2.4	0.40

Description after deal with outliers by IQR:

```
dataset.describe()
```

	Age	Gender	Total_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	Albumin	Albumin_and_Globulin_Ratio
count	583.000000	583.000000	583.000000	583.000000	583.000000	583.000000	583.000000	583.000000	583.000000
mean	44.746141	0.756432	2.249400	266.389365	53.399657	73.041166	6.483190	3.141852	0.945403
std	16.189833	0.429603	2.382344	145.665460	46.059536	73.549864	1.085451	0.795519	0.310942
min	4.000000	0.000000	0.400000	63.000000	10.000000	10.000000	2.700000	0.900000	0.300000

25%	33.000000 5.800000	1.000000 2.600000	0.800000 0.700000	175.500000 1.000000	23.000000	25.000000
50%	45.000000 6.600000	1.000000 3.100000	1.000000 0.930000	208.000000 1.000000	35.000000	42.000000
75%	58.000000 7.200000	1.000000 3.800000	2.600000 1.100000	298.000000 2.000000	60.500000	87.000000
max	90.000000 9.600000	1.000000 5.500000	8.000000 2.300000	665.500000 2.000000	173.000000	273.000000

Independent and Dependent Split

Independent and Dependent Feature:

X = dataset.iloc[:, :-1]

y = dataset.iloc[:, -1]

top 5 records of Independent features:

X.head()

Age	Gender	Total_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	Albumin	Albumin_and_Globulin_Ratio	
0	65	0	0.7	187.0	16	18	6.8	3.3	0.90
1	62	1	8.0	665.5	64	100	7.5	3.2	0.74
2	62	1	7.3	490.0	60	68	7.0	3.3	0.89
3	58	1	1.0	182.0	14	20	6.8	3.4	1.00
4	72	1	3.9	195.0	27	59	7.3	2.4	0.40

top 5 records of dependent features:

y.head()

0 1

1 1

2 1

3 1

4 1

Name: Dataset, dtype: int64

SMOTE Technique:

```
from imblearn.combine import SMOTETomek
```

```
smote = SMOTETomek()
```

```
X_smote, y_smote = smote.fit_resample(X,y)
```

Counting before and after SMOTE:

```
from collections import Counter
```

```
print('Before SMOTE : ', Counter(y))
```

```
print('After SMOTE : ', Counter(y_smote))
```

Before SMOTE : Counter({1: 416, 2: 167})

After SMOTE : Counter({1: 396, 2: 396})

Train Test Split

Train Test Split:

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(X_smote,y_smote, test_size=0.3, random_state=33)
```

```
print(X_train.shape)
```

```
print(X_test.shape)
```

(554, 9)

(238, 9)

Feature Importance :

```
from sklearn.feature_selection import SelectKBest
```

```
from sklearn.feature_selection import chi2
```

```
### Apply SelectKBest Algorithm
```

```
ordered_rank_features=SelectKBest(score_func=chi2,k=9)
```

```
ordered_feature=ordered_rank_features.fit(X,y)
```

```
dfscores=pd.DataFrame(ordered_feature.scores_,columns=["Score"])
```

```
dfcolumns=pd.DataFrame(X.columns)
```

```
features_rank=pd.concat([dfcolumns,dfscores],axis=1)
```

```
features_rank.columns=['Features','Score']
```

```
features_rank.nlargest(9, 'Score')
```

	Features	Score
5	Aspartate_Aminotransferase	3368.743077
3	Alkaline_Phosphotase	2385.790640
4	Alamine_Aminotransferase	1717.348297
2	Total_Bilirubin	127.476411
0	Age	64.315174
7	Albumin	3.053371
8	Albumin_and_Globulin_Ratio	1.704602
1	Gender	0.964518
6	Total_Protiens	0.129627

```
Classification Algorithm And Run model
```

```
# Importing Performance Metrics:
```



```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
# RandomForestClassifier:
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
RandomForest = RandomForestClassifier()
```

```
RandomForest = RandomForest.fit(X_train,y_train)
```

```
# Predictions:
```

```
y_pred = RandomForest.predict(X_test)
```

```
# Performance:
```

```
print('Accuracy:', accuracy_score(y_test,y_pred))
```

```
print(confusion_matrix(y_test,y_pred))
```

```
print(classification_report(y_test,y_pred))
```

```
Accuracy: 0.8109243697478992
```

```
[[ 87 35]
```

```
 [ 10 106]]
```

```
precision recall f1-score support
```

```
1    0.90    0.71    0.79    122
```

```
2    0.75    0.91    0.82    116
```

```
accuracy                0.81    238
```

```
macro avg    0.82    0.81    0.81    238
```

```
weighted avg    0.83    0.81    0.81    238
```

```
# AdaBoostClassifier:

from sklearn.ensemble import AdaBoostClassifier

AdaBoost = AdaBoostClassifier()

AdaBoost = AdaBoost.fit(X_train,y_train)
```

```
# Predictions:

y_pred = AdaBoost.predict(X_test)
```

```
# Performance:

print('Accuracy:', accuracy_score(y_test,y_pred))

print(confusion_matrix(y_test,y_pred))

print(classification_report(y_test,y_pred))
```

```
Accuracy: 0.7857142857142857
```

```
[[ 86 36]
```

```
 [ 15 101]]
```

```
precision recall f1-score support
```

```
1    0.85    0.70    0.77    122
```

```
2    0.74    0.87    0.80    116
```

```
accuracy                0.79    238
```

```
macro avg    0.79    0.79    0.78    238
```

```
weighted avg    0.80    0.79    0.78    238
```

```
# GradientBoostingClassifier:
```

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
GradientBoost = GradientBoostingClassifier()
```

```
GradientBoost = GradientBoost.fit(X_train,y_train)
```

```
# Predictions:
```

```
y_pred = GradientBoost.predict(X_test)
```

```
# Performance:
```

```
print('Accuracy:', accuracy_score(y_test,y_pred))
```

```
print(confusion_matrix(y_test,y_pred))
```

```
print(classification_report(y_test,y_pred))
```

```
Accuracy: 0.7605042016806722
```

```
[[89 33]
```

```
[24 92]]
```

```
precision recall f1-score support
```

```
1    0.79    0.73    0.76    122
```

```
2    0.74    0.79    0.76    116
```

```
accuracy                0.76    238
```

```
macro avg    0.76    0.76    0.76    238
```

```
weighted avg    0.76    0.76    0.76    238
```

```
RandomizedSearchCV
```

```
# Importing RandomizedSearchCV:
```

```
from sklearn.model_selection import RandomizedSearchCV

# Number of trees in random forest:

n_estimators = [int(x) for x in np.linspace(start = 100, stop = 2000, num = 20)]

# Number of features to consider at every split:

max_features = ['auto', 'sqrt', 'log2']

# Maximum number of levels in tree:

max_depth = [int(x) for x in np.linspace(100, 100, 20)]

# Minimum number of samples required to split a node:

min_samples_split = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 16, 18, 20]

# Minimum number of samples required at each leaf node:

min_samples_leaf = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 16, 18, 20]

# Create the random grid:

random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'criterion':['entropy','gini']}

print(random_grid)
```

```
4, 5, 6, 7, 8, 9, 10, 12, 14, 16, 18, 20], 'min_samples_leaf': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 16, 18, 20],  
'criterion': ['entropy', 'gini']}]
```

```
rf = RandomForestClassifier()
```

```
rf_randomcv = RandomizedSearchCV(estimator = rf, param_distributions = random_grid, n_iter = 100,  
cv = 5, verbose = 2,
```

```
random_state = 0, n_jobs = -1)
```

```
# fit the randomized model:
```

```
rf_randomcv.fit(X_train,y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py:372: FitFailedWarning:

40 fits failed out of a total of 500.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

40 fits failed with the following error:

Traceback (most recent call last):

File "/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py", line 680, in
_fit_and_score

```
estimator.fit(X_train, y_train, **fit_params)
```

File "/usr/local/lib/python3.7/dist-packages/sklearn/ensemble/_forest.py", line 467, in fit

```
for i, t in enumerate(trees)
```

File "/usr/local/lib/python3.7/dist-packages/joblib/parallel.py", line 1085, in __call__

```
if self.dispatch_one_batch(iterator):
```

File "/usr/local/lib/python3.7/dist-packages/joblib/parallel.py", line 901, in dispatch_one_batch

```
self._dispatch(tasks)
```

```
File "/usr/local/lib/python3.7/dist-packages/joblib/parallel.py", line 819, in _dispatch
```

```
    job = self._backend.apply_async(batch, callback=cb)
```

```
File "/usr/local/lib/python3.7/dist-packages/joblib/_parallel_backends.py", line 208, in apply_async
```

```
    result = ImmediateResult(func)
```

```
File "/usr/local/lib/python3.7/dist-packages/joblib/_parallel_backends.py", line 597, in __init__
```

```
    self.results = batch()
```

```
File "/usr/local/lib/python3.7/dist-packages/joblib/parallel.py", line 289, in __call__
```

```
    for func, args, kwargs in self.items]
```

```
File "/usr/local/lib/python3.7/dist-packages/joblib/parallel.py", line 289, in
```

```
    for func, args, kwargs in self.items]
```

```
File "/usr/local/lib/python3.7/dist-packages/sklearn/utils/fixes.py", line 216, in __call__
```

```
    return self.function(*args, **kwargs)
```

```
File "/usr/local/lib/python3.7/dist-packages/sklearn/ensemble/_forest.py", line 185, in  
_parallel_build_trees
```

```
    tree.fit(X, y, sample_weight=curr_sample_weight, check_input=False)
```

```
File "/usr/local/lib/python3.7/dist-packages/sklearn/tree/_classes.py", line 942, in fit
```

```
    X_idx_sorted=X_idx_sorted,
```

```
File "/usr/local/lib/python3.7/dist-packages/sklearn/tree/_classes.py", line 254, in fit
```

```
    % self.min_samples_split
```

```
ValueError: min_samples_split must be an integer greater than 1 or a float in (0.0, 1.0]; got the integer 1
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_search.py:972: UserWarning: One or  
more of the test scores are non-finite: [    nan 0.6950041  0.74733825 0.68958231 0.68599509  
0.71847666
```

```
0.73829648 0.7040131 0.73290745 0.72029484 0.69320229 0.72209664
```

0.6950041 0.69141687 0.6986077 0.73113841 0.68781327 0.74733825
0.72211302 0.75094185 nan 0.71665848 0.74194922 0.71667486
0.70044226 0.69678952 0.72026208 0.71665848 0.74552007 0.71846028
0.69138411 nan 0.72208026 0.72568387 0.69316953 0.72568387
0.72027846 0.75099099 0.74552007 0.73467649 0.71665848 0.69140049
0.69682228 0.70042588 0.68782965 0.73112203 nan 0.70584767
nan 0.72748567 0.71667486 0.6968059 0.70763309 0.72750205
0.71847666 0.74735463 0.75276003 0.73290745 0.6968059 0.72748567
0.71485667 0.71841114 0.68781327 0.72031122 0.75276003 0.6968059
0.73289107 0.72208026 0.73108927 0.7004095 nan 0.73829648
0.7040131 0.72751843 0.73651106 0.72751843 0.72568387 0.72209664
0.70583129 0.68963145 0.73651106 0.69502048 0.72209664 nan
0.68959869 0.74915643 0.72930385 0.74917281 0.73108927 0.74191646
0.74191646 0.74373464 nan 0.73290745 0.6968059 0.68959869
0.71667486 0.69498771 0.73832924 0.71305487]

category=UserWarning,

RandomizedSearchCV(cv=5, estimator=RandomForestClassifier(), n_iter=100,

n_jobs=-1,

param_distributions={'criterion': ['entropy', 'gini'],

'max_depth': [100, 100, 100, 100, 100,

100, 100, 100, 100, 100,

100, 100, 100, 100, 100,

100, 100, 100, 100, 100],

'max_features': ['auto', 'sqrt',

'log2'],

```

        'min_samples_leaf': [1, 2, 3, 4, 5, 6,
                               7, 8, 9, 10, 12,
                               14, 16, 18, 20],
        'min_samples_split': [1, 2, 3, 4, 5, 6,
                               7, 8, 9, 10, 12,
                               14, 16, 18, 20],
        'n_estimators': [100, 200, 300, 400,
                          500, 600, 700, 800,
                          900, 1000, 1100, 1200,
                          1300, 1400, 1500, 1600,
                          1700, 1800, 1900,
                          2000]},
        random_state=0, verbose=2)

# Best parameter of RandomizedSearchCV:

rf_randomcv.best_params_

{'n_estimators': 2000,
 'min_samples_split': 3,
 'min_samples_leaf': 2,
 'max_features': 'log2',
 'max_depth': 100,
 'criterion': 'entropy'}

# Creating model using best parameter of RandomizedSearchCV:

RandomForest_RandomCV = RandomForestClassifier(criterion = 'entropy', n_estimators = 2000,
max_depth = 100, max_features = 'log2',

        min_samples_split = 3, min_samples_leaf = 2)

RandomForest_RandomCV = RandomForest_RandomCV.fit(X_train,y_train)

```


Predictions:

```
y_pred = RandomForest_RandomCV.predict(X_test)
```

Performance:

```
print('Accuracy:', accuracy_score(y_test,y_pred))
```

```
print(confusion_matrix(y_test,y_pred))
```

```
print(classification_report(y_test,y_pred))
```

Accuracy: 0.8067226890756303

[[86 36]

[10 106]]

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

1	0.90	0.70	0.79	122
---	------	------	------	-----

2	0.75	0.91	0.82	116
---	------	------	------	-----

accuracy			0.81	238
----------	--	--	------	-----

macro avg	0.82	0.81	0.81	238
-----------	------	------	------	-----

weighted avg	0.82	0.81	0.80	238
--------------	------	------	------	-----

GridSearchCV

Importing GridSearchCV:

```
from sklearn.model_selection import GridSearchCV
```

Best parameter:

```
rf_randomcv.best_params_
```

```

{'n_estimators': 2000,
 'min_samples_split': 3,
 'min_samples_leaf': 2,
 'max_features': 'log2',
 'max_depth': 100,
 'criterion': 'entropy'}

param_grid = {
    'criterion': [rf_randomcv.best_params_['criterion']],
    'max_features': [rf_randomcv.best_params_['max_features']],
    'max_depth': [rf_randomcv.best_params_['max_depth']-50,
                  rf_randomcv.best_params_['max_depth'],
                  rf_randomcv.best_params_['max_depth']+50],
    'min_samples_leaf': [rf_randomcv.best_params_['min_samples_leaf']-1,
                         rf_randomcv.best_params_['min_samples_leaf'],
                         rf_randomcv.best_params_['min_samples_leaf']+1],
    'min_samples_split': [rf_randomcv.best_params_['min_samples_split'] - 1,
                          rf_randomcv.best_params_['min_samples_split'],
                          rf_randomcv.best_params_['min_samples_split'] + 1],
    'n_estimators': [rf_randomcv.best_params_['n_estimators'] - 50,
                     rf_randomcv.best_params_['n_estimators'],
                     rf_randomcv.best_params_['n_estimators'] + 50]
}

print(param_grid)

{'criterion': ['entropy'], 'max_features': ['log2'], 'max_depth': [50, 100, 150], 'min_samples_leaf': [1, 2, 3],
 'min_samples_split': [2, 3, 4], 'n_estimators': [1950, 2000, 2050]}

```

Fit the grid_search to the data:

```
rf = RandomForestClassifier()
```

```
grid_search = GridSearchCV(estimator = rf, param_grid = param_grid, cv=5 , n_jobs = -1, verbose = 2)
```

```
grid_search.fit(X_train,y_train)
```

Fitting 5 folds for each of 81 candidates, totalling 405 fits

```
GridSearchCV(cv=5, estimator=RandomForestClassifier(), n_jobs=-1,  
             param_grid={'criterion': ['entropy'], 'max_depth': [50, 100, 150],  
                         'max_features': ['log2'],  
                         'min_samples_leaf': [1, 2, 3],  
                         'min_samples_split': [2, 3, 4],  
                         'n_estimators': [1950, 2000, 2050]},  
             verbose=2)
```

Best Parameter of GridSearchCV:

```
grid_search.best_params_
```

```
{'criterion': 'entropy',
```

```
  'max_depth': 50,
```

```
  'max_features': 'log2',
```

```
  'min_samples_leaf': 2,
```

```
  'min_samples_split': 2,
```

```
  'n_estimators': 2050}
```

Creating model using best parameter of GridSearchCV:

```
RandomForest_gridCV = RandomForestClassifier(criterion='entropy', n_estimators=1950,  
max_depth=150, max_features='log2',
```

```
      min_samples_split=2, min_samples_leaf=1)
```

```
RandomForest_gridCv = RandomForest_gridCV.fit(X_train,y_train)
```

Predictions:

```
y_pred = RandomForest_gridCV.predict(X_test)
```

Performance:

```
print('Accuracy:', accuracy_score(y_test,y_pred))
```

```
print(confusion_matrix(y_test,y_pred))
```

```
print(classification_report(y_test,y_pred))
```

Accuracy: 0.8025210084033614

```
[[ 85 37]
```

```
 [ 10 106]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

1	0.89	0.70	0.78	122
---	------	------	------	-----

2	0.74	0.91	0.82	116
---	------	------	------	-----

accuracy			0.80	238
----------	--	--	------	-----

macro avg	0.82	0.81	0.80	238
-----------	------	------	------	-----

weighted avg	0.82	0.80	0.80	238
--------------	------	------	------	-----

Save Model

Creating a pickle file for the classifier

```
import pickle
```

```
filename = 'Liver.pkl'
```

```
pickle.dump(RandomForestClassifier, open(filename, 'wb'))
```